

# DAVID JEREMIE ANAND

Applied AI Engineer | Agentic Systems & RAG | LLM Evaluation

New Delhi, India | +91 70118 90891 | jerryanand10@gmail.com

davidjeremieanand.com | linkedin.com/in/david-jeremie-anand | github.com/Jeremie2002-sudo

## PROFILE

Builds AI systems, then proves they work: **four shipped AI projects** across agentic triage, RAG, tool calling, and LLM evaluation – each with its own **eval harness and CI suite**, three live on Render. Two years in Genpact operations supply the instincts those systems automate: SLA pressure, escalation judgment, root-cause discipline. Backed by a PGP from Masters' Union.

## PROFESSIONAL EXPERIENCE

### Genpact

Aug 2024 – Aug 2026

New Delhi, India

*Associate and Subject Matter Expert – Operations*

- Analysed workflows handling **100,000+ cases a year** and drove process redesign, contributing to an **18% reduction in turnaround time**.
- Ran structured root-cause analysis of recurring defects and SLA deviations, tracking issues through Buganizer and validating corrective actions, contributing to a **30% reduction in operational defects** with minimal SLA breaches.
- Automated management reports and dashboards in Excel and Power BI, cutting recurring reporting effort by **25%** and saving about **100 FTE hours annually**.
- Partnered with Risk, Compliance, and Quality stakeholders on controls and workflow improvements, contributing to a **22% productivity uplift** across a ~20-person operations team.

## TECHNICAL PROJECTS

### FirstLine – AI Ops Triage Agent with Self-Scoring Evals [GitHub](#) · [Live](#)

Python, FastAPI, Ollama/Groq, pytest, GitHub Actions

- Built an agentic first line of support that works a full ops shift – classifying root causes across 15 incident archetypes, clustering cross-service incidents, tracking SLA clocks, escalating judgment calls, and drafting RCAs – over a seeded simulator with hidden labeled ground truth.
- Designed an eval harness that grades the agent against that truth (triage accuracy, missed POs, escalation precision); live Llama 3.2 runs surfaced **100% escalation recall at 21% precision** – safety by crying wolf; shipped with a **768-test CI suite** and a live triage dashboard.

### Synapse – Agentic RAG Document Assistant [GitHub](#) · [Live](#)

Python, FastAPI, FAISS, Ollama/Groq, pytest, GitHub Actions

- Built a hybrid retrieval pipeline (BM25 + FAISS fused via Reciprocal Rank Fusion, LLM listwise reranking) feeding a 4-stage agent pipeline (planner, retriever, synthesizer, critic) that grounds answers in evidence and self-critiques unsupported claims.
- Added an observability UI and an automated eval harness for retrieval and faithfulness scoring, plus a pluggable model backend (local Ollama / hosted Groq); shipped with a **31-test CI suite**, deployed on Render.

### Ledgr – AI-Powered Finance Manager [GitHub](#) · [Live](#)

Python, Flask, SQLite, OpenAI/Ollama, Docker

- Built a full-stack finance tracker with dashboards, budget tracking, and CSV import/export, plus AI transaction categorisation (concurrent LLM calls with keyword fallback) that works fully offline.
- Designed a tool-calling assistant that queries real transaction data instead of relying on LLM arithmetic, with a pluggable OpenAI/Ollama backend; shipped with a **58-test CI suite**, Dockerised, deployed via Render Blueprint.

### PromptDrift – LLM Prompt Regression Testing [GitHub](#)

Python, asyncio, Ollama/Groq, pytest, GitHub Actions

- Built a CLI that snapshots prompt behaviour into a committed baseline and fails CI when a change degrades it, sampling every case repeatedly so genuine regressions are separated from model non-determinism.
- Tiered assertions run deterministic checks before paid LLM-as-judge calls, cutting evaluation cost; shipped with a **157-test CI suite** across Python 3.10–3.13, running fully offline through a deterministic mock backend.

### Wireless Temperature Detection System using HackRF One

RF and Wireless Systems

- Built an RF prototype on HackRF One implementing the end-to-end SDR chain (capture, filtering, demodulation, decoding) to read transmitted temperature signals, validated against a reference sensor.

## INTERNSHIPS

### Bhagwan Parshuram Institute of Technology, GGSIPU

Jul 2023 – Aug 2023

*Data Analytics Intern*

New Delhi, India [Certificate](#)

- Analysed institutional datasets in Python and Excel and built KPI dashboards and automated reports, cutting repetitive manual effort by **25%**.

### Defence Research and Development Organisation (DRDO)

Jul 2022 – Sep 2022

*Research and Development Intern*

India [Certificate](#)

- Supported development, validation, and signal analysis of an embedded wireless sensing prototype, improving reliability and reducing errors.

### AGIXURY

Sep 2020 – Oct 2020

*Web Development Intern*

Remote [Certificate](#)

- Developed responsive web interfaces from business requirements and improved usability across desktop and mobile layouts.

## EDUCATION

### Masters' Union

Aug 2026 – Present

*PGP in Applied AI and Agentic Systems*

Gurugram, India

### Bhagwan Parshuram Institute of Technology, GGSIPU

2020 – 2024

*B.Tech in Electronics and Communication Engineering* | CGPA: 8.53/10.0

New Delhi, India

*Faith Academy Senior Secondary School, CBSE* Class XII: 92.8%, Class X: 94.2% | New Delhi, India

2018 – 2020

## TECHNICAL SKILLS

|                                 |                                                                                                                                                                                                                                                                                                             |
|---------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Programming and Data            | Python, SQL, SQLite, FastAPI, Flask, REST APIs, NumPy, asyncio, data analysis, exploratory data analysis                                                                                                                                                                                                    |
| Applied AI                      | Machine learning, generative AI, agentic AI, multi-agent pipelines, large language models, embeddings, prompt engineering, tool calling (function calling), Retrieval-Augmented Generation (RAG)                                                                                                            |
| Retrieval, Evaluation and Cloud | Hybrid retrieval (BM25 + FAISS), reciprocal rank fusion, LLM reranking, vector databases, LLM evaluation (LLM-as-a-judge), faithfulness scoring, eval harness design, prompt regression testing, observability, LLM Ops, GenAIOps, model APIs (OpenAI, Ollama, Groq), Microsoft Azure AI, Snowflake, Render |
| Analytics and Development Tools | Power BI, Tableau, Advanced Excel, Git, GitHub Actions, CI/CD, Docker, pytest, VS Code                                                                                                                                                                                                                      |
| Delivery and Problem Solving    | Root-cause analysis, stakeholder communication, process improvement, customer-facing delivery                                                                                                                                                                                                               |

## SELECTED CERTIFICATIONS

Professional Certificates: [Snowflake Generative AI](#) | [Microsoft Azure AI Essentials](#) | [Astronomer Data Engineering Foundations](#)

AI Engineering: [Become an AI Engineer](#) | [GenAIOps Foundations](#) | [Generative AI in Cloud Computing](#)

LLMs and RAG: [Fine-Tuning for LLMs](#) | [Advanced LLMs with RAG](#) | [Vector Databases and RAG](#) | [AI Ecosystem for Developers](#)